

# EMBEDDED DEVICE DRIVERS

---

Linux Device Drivers on Beaglebone Black

# LKM: Device Files

- This course focus:
  - Character devices
- So far
  - We assigned ***major:minor*** numbers
  - But we lack the device “file”
    - So user space can not interact with our device
    - Neither can our device interact with any hardware
- Let's create a device node / file entry in ***/dev***

# LKM: Device node creation

- Device files / nodes enable interaction
  - Between user space and hardware
  - Not normal files, and usually created in **/dev**
- Once created, user space programs can
  - open, read, write, close, mmap(), etc.
- When user space program requests an operation
  - The kernel traps this operation
  - And passes it to our device driver
  - Which then performs the operation
  - On behalf of the user space program
- Device nodes can be created **manually** / **programmatically**

# LKM: Node creation (manual)

- Shell command

`$ mknod -m <perms> <dev_name> <dev_type> <major> <minor>`

- *perms*: permissions on the device
- *dev\_name*: Full path name (/dev/...)
- *dev\_type*: **c** for char, **b** for block
- *major, minor*: Numbers assigned to this device

- Example:

`$ sudo mknod -m 0755 /dev/my_device c 202 0`

- Note that device node can be created

- Only by the root user
- Even before loading the device driver!

# LKM: Node creation (auto)

- Programmatic creation of a dev node
  - Follows a hierarchy:
    - Class → Device
  - Actions:
    - Create a class
      - Created under **/sys/class**
    - Create a device within this class
      - Created in **sysfs**
- Header files:
  - `#include <linux/device.h>`
  - `#include <linux/kdev_t.h>`

# LKM: Node creation API

- Class creation

*struct class \***class\_create**(struct module \*owner, const char \*name);*

- *owner*: Who owns this class (*THIS\_MODULE*)
- *name*: String name for this class

- Device creation

*struct device \***device\_create**(struct \*class, struct device \*parent, dev\_t dev, void \*drvdata, const char \*fmt, ...);*

- *class*: Class under which device is registered
- *parent*: Parent device in hierarchy
- *dev*: The *dev\_t* for the char device
- *drvdata*: Data to be added to the device for callbacks
- *fmt*: String of the device name

# LKM: Node cleanup API

- Class destroy

*void **class\_destroy**(struct class \*cls);*

- Device destroy

*void **device\_destroy**(struct class \*cls, dev\_t dev);*

# LKM: Device node exercise

- Refer ***mod3*** directory
  - File ***mod33.c*** contains the src code
    - We seek allocation for a *major:minor* from kernel
    - We create a device class (***cdac\_cls***)
    - We create a char device under this class (***cdac\_dev***)
    - Compile, load the module
    - Observe the dmesg output to get ***major:minor*** allocation
    - Also check the entries created:
      - `$ ls -l /sys/class | grep cdac`
      - `$ ls -l /sys/dev/char/ | grep cdac`
      - `$ ls -l /dev | grep cdac`
    - Unload the module



# LKM: File Operations

- So far, for our char device, we have managed to
  - Assign **major:minor** (static / dynamic)
  - Create class and device nodes
  - But how do we use this device?
    - Through file operations
    - Since the device node is basically a “file” under Linux
- To perform file operations
  - (**fops** in kernel-speak)
  - We need to register structures
    - That tell the kernel what to do
    - When a particular file operation needs to be performed
  - Most driver operations use:
    - **struct cdev, struct file\_operations, struct inode**

# LKM: struct cdev

- Internal structure
  - Used by kernel
  - For char devices

- Relevant elements
  - *owner*
    - *THIS\_MODULE*
  - *ops*
    - For file operations

```
#include <linux/kdev_t.h>
```

```
struct cdev {
```

```
...
```

```
struct module *owner;
```

```
const struct  
file_operations *ops;
```

```
...
```

```
};
```

# LKM: struct cdev API

- Allocation
  - Embedded  
`void cdev_init(struct cdev *cdev, struct file_operations *fops);`
  - Runtime  
`struct cdev *my_cdev = cdev_alloc();`  
`my_cdev->ops = &my_fops;`
- Registering with kernel  
`int cdev_add(struct cdev *cdev, dev_t dev, unsigned int count);`
  - `cdev`: cdev structure
  - `dev`: major:minor holder
  - `count`: No of devices to associate with the device
- De-registering from the kernel  
`void cdev_del(struct cdev *cdev);`

# LKM: struct file\_operations

```
#include <linux/fs.h>

struct file_operations {
    struct module *owner;
    loff_t (*llseek) (struct file *, loff_t, int);
    ssize_t (*read) (struct file *, char __user *, size_t,
loff_t *);
    ssize_t (*write) (struct file *, const char __user *,
size_t, loff_t *);
    [...]
    long (*unlocked_ioctl) (struct file *, unsigned int,
unsigned long);
    [...]
    int (*open) (struct inode *, struct file *);
    int (*flush) (struct file *, fl_owner_t id);
    int (*release) (struct inode *, struct file *);
    [...]
```

- Has function pointers
  - For registering various operations
  - On char devices
- Notice APIs differ from system calls
  - Kernel mediates
    - Between user and device driver
    - Simplifies user space calls

# LKM: File perspectives

- We have 2 structs for files
  - ***struct inode***
    - Represents a file from the file-system viewpoint
    - Attributes: size, rights, access/modification times
    - Uniquely identifies a file in a file-system
    - *Used to get the major:minor of the device*
  - ***struct file***
    - Represents a file from the user's viewpoint
    - Attributes: inode, file-name, opening options, position
    - All open files have an associated *struct file*
    - *Used to determine flags and for operations on the file*

# LKM: On file\_operations (1/2)

- Owner
  - ***THIS\_MODULE***
  - Tells the kernel not to unload the module when an operation is in progress
- Open
  - int (\***open**)(struct inode \*, struct file\*);*
  - Usually, the first operation performed on the device
    - But not mandatory – in which case device open always succeeds
- Read
  - ssize\_t (\***read**)(struct file \*, char \_\_user \*, size\_t, loff\_t \*);*
  - Used to retrieve data from the device
    - Non-negative return value indicates bytes read
    - Negative return value indicates error

# LKM: On file\_operations (2/2)

- Write

*ssize\_t (\***write**)(struct file \*, const char \_\_user \*, size\_t, loff\_t \*);*

- Used to send data to the device
  - Non-negative return value indicates bytes written
  - Negative value indicates error

- Release (close)

*int (\***release**)(struct inode \*, struct file \*);*

- Usually called when device is closed

- IOCTL

*long (\***ioctl**)(struct file \*, unsigned int, unsigned long);*

- Used to issue device specific commands that are not data specific

# LKM: fops exercise (1/2)

- Refer **mod3** directory
  - File **mode34.c** contains the src code
    - Here, we create a char device (**/dev/cdac\_dev**) like in **mod33**
    - We create a **fops** using our own file operations
    - We create a **cdev** and register it with the kernel using the **fops**
  - Compile and load the module on BBB
  - Once the module has loaded
    - Test the read part
      - # cat /dev/cdac\_dev**
        - Observe *dmesg*
    - Test the write part
      - # echo 123 > /dev/cdac\_dev**
        - Observe *dmesg*



# LKM: fops exercise (2/2)

- Refer ***mod3*** directory
  - File ***mode34\_app.c*** contains the src code
    - User space code to test the char device driver
      - We open the device (***/dev/cdac\_dev***)
      - Do a write
      - Do a read
  - All using system calls
  - Compile the .c file on BBB
  - Notice the output of ***dmesg*** when you run ***a.out***
    - When module is unload
    - After module is loaded

# LKM: Kernel memory allocation

- Kernel has its own memory space
  - Shared with scheduler, device drivers, etc.
- Kernel manages its own memory space
  - Using kernel-level memory allocators
  - Cannot use C library `<stdlib.h>` (*malloc*, *calloc*, *free*)
- Kernel memory management APIs
  - *kmalloc()*
  - *kfree()*

# LKM: Memory mgmt. APIs

- Allocation

`void *kmalloc(size_t size, gfp_t flags);`

- `size`: Bytes of memory requested
- `flags`: Type of memory needed
  - `GFP_USER`: Alloc on behalf of user, may sleep
  - **`GFP_KERNEL`**: Alloc kernel RAM, may sleep
  - `GFP_ATOMIC`: No sleep, use emergency pools
  - `GFP_NOIO`: No IO while allocating memory
  - `__GFP_THISNODE`: Alloc on this node only
  - `GFP_DMA`: Alloc memory to used for DMA

- Freeing

`void kfree(const void *objp);`

- `objp`: pointer returned by `kmalloc()`

- Header file

**`#include <linux.slab.h>`**

# LKM: Interactions – user/kernel

- In many applications
  - User space programs
  - Need to interact with
  - Kernel space code (like modules / device drivers)
  - Data needs to be exchanged between the spaces
    - From user to kernel
    - From kernel to user
- Linux provides specific APIs for these interactions
  - Header file
    - #include <linux/uaccess.h>***

# LKM: Userspace → kernelspace

- Copy memory from user space to kernel space  
*unsigned long **copy\_from\_user**(void \*to, const void \_\_user \*from, unsigned long n);*
  - *to*: Destination in kernel space
  - *from*: Source in user space
  - *n*: Bytes to copy
- Returns:
  - No of bytes that **could not be copied**
  - So success is 0!

# LKM: Kernel space → Userspace

- Copy memory from kernel space to user space  
*unsigned long **copy\_to\_user**(void \_\_user \*to, const void \*from, unsigned long n);*
  - *to*: Destination in user space
  - *from*: Source in kernel space
  - *n*: Bytes to copy
- Returns:
  - No of bytes that **could not be copied**
  - So success is 0!

# LKM: Interaction b/w spaces

- We now look at passing data
  - From user space to kernel space
  - And vice versa
- This needs 2 pieces of code
  - Device driver
    - To emulate a dummy device
    - Accepts data from user space and passes it back
  - User space program
    - Regular code using systems calls
    - Interacts with device driver

# LKM: Space interaction exercise

- Refer ***mod3*** directory
  - The file ***mod35.c*** is the driver src code
    - Creates a character device as before
    - Read is modified to accept user space buffer
    - Write is modified to write to user space buffer
    - Compile and load the module on BBB
  - The file ***mod35\_app.c*** is the user space program
    - We open the device (***/dev/cdac\_dev***)
    - We first write to it; then we read from it
    - Compile the .c file on BBB to get ***a.out***
    - Observe the ***dmesg*** output when you run ***a.out***
      - With module loaded
      - Without module loaded



THANK YOU!