

# EMBEDDED DEVICE DRIVERS

---

Linux Device Drivers on Beaglebone Black

# LKM: Kernel Threads

- Like user space, Linux kernel supported threading
  - Threads are called kernel threads (kthreads)
  - Run purely in kernel space
    - Hence, user space thread management framework **does not** apply
      - Such as pthreads, etc.
- Kernel thread usage steps
  - Creation
  - Start
  - Stop – by others / self
  - Other functions – waiting, completion, etc.
- Header file:  
***#include <linux/kthread.h>***

# LKM: kthread creation/stop

- Creation - thread has to be “woken up” to start running

*struct task\_struct \***kthread\_create**(int (\*thread\_func)(void \*data), const char namefmt[], ...);*

- *thread\_func*: Function to run in the created thread; API: **int (\*func)(void \*)**;
- *namefmt*: Printf style name for the thread

- Returns:

- **task\_struct** \* on success
- **ERR\_PTR(-ENOMEM)** on failure

- Wake up

*int **wake\_up\_process**(struct task\_struct \*thread);*

- Combined create-and-wake-up/run wrapper

*struct task\_struct \***kthread\_run**(int (\*thread\_func)(void \*data), void \*data, const char namefmt[]);*

- Stop

*int **kthread\_stop**(struct task\_struct \*thread);*

- Sets the **kthread\_stop** variable of **thread**

# LKM: kthread self-exit

- Each kthread has an internal variable
  - ***kthread\_stop***
    - Set/Reset by somebody other than this kthread (ideally!)
- Long-running kthread should
  - Check value of its ***kthread\_stop*** variable
    - ***kthread\_should\_stop()***
      - If non-zero, it should return
- A kthread can end itself; optionally returning a value  
`void do_exit(long retcode);`
  - Used in “detached” mode (from ***kthread*** parent)

# LKM: Completion (kthread)

- Linux kernel has a mechanism
  - To indicate whether some work is done / thread can self-exit
- Called completion
  - Header: ***#include <linux/completion.h>***
  - Variable: *struct **completion** my\_cmpln;*
  - APIs:
    - Creating:
      - Static: ***DECLARE\_COMPLETION***(my\_cmpln);
      - Dynamic: ***init\_completion***(struct completion \*my\_cmpln);
    - Waiting:
      - void **wait\_for\_completion**(struct completion \*my\_cmpln);*
    - Completing:
      - void **complete**(struct completion \*my\_cmpln);*
    - Checking status:
      - bool **completion\_done**(struct completion \*my\_cmpln);*

# LKM: kthread Exercise #1

- Refer **mod8** directory
  - **mod81.c** contains code for basic kthreads
    - Study the way threads are created in the kernel
      - And how it differs from user space!
    - Compile and load the module on BBB
    - Observe **dmesg** output
  - **mod82.c** contains code for kthreads with completion
    - Study how a wait thread can be created
      - And made to wait for completion
      - And made to self-exit upon getting the completion
    - Compile and load the module on BBB
    - Test read with **cat /dev/cdac\_dev**
      - Observe **dmesg** output
    - Unload the module
      - Observe **dmesg** output

# LKM: Sharing b/w kthreads

- We know things go downhill
  - When resources are shared
  - Between threads
  - Without proper synchronization
- kthreads suffer from the same malaise
  - Hence we need thread-level synchronization primitives
    - We already know (from user space)
      - Semaphores
      - **Mutexes**
    - Kernel has some more
      - **Spinlocks**
      - Read-write locks
      - Seqlocks
      - Atomic variables

# LKM: kthread Exercise #2

- Refer ***mod8*** directory
  - ***mod83.c*** contains code for unprotected shared resources
    - We share an integer and a char buffer b/w 2 kthreads
    - Compile and load the module on BBB
    - Observe ***dmesg*** output
      - What do you see as far as kthread-1's prints are concerned?
  - ***mod84.c*** contains code for protection using a mutex
    - Study how we create a mutex, init it, and perform locks/unlocks
    - Compile and load the module on BBB
    - Observe ***dmesg*** output
      - What do you observe in terms of kthread prints?



# LKM: Spinlocks

- In the mutex case
  - When a kthread is trying to acquire a lock
    - If the mutex is not available
      - It goes to sleep
      - Wakes when it gets the mutex lock
- Spinlocks are similar to mutex
  - Single-owner locks
  - Locks / unlock operations
  - But the kthread wanting to lock a spinlock
    - **Spins** – busy spin of the underlying CPU
    - Till it gets the lock
    - No sleep induced
  - Advantage: A fast and simple locking mechanism
  - Restrictions: Waits have to be really short!

# LKM: Spinlock API (simple)

- Creation

- Static

- ```
DEFINE_SPINLOCK(my_spinlock);
```

- Dynamic

- ```
spinlock_t my_spinlock;
```

- ```
spin_lock_init(&my_spinlock);
```

- Locking

- ```
spin_lock(spinlock_t *lock);
```

- ```
spin_trylock(spinlock_t *lock);
```

- Unlocking

- ```
spin_unlock(spinlock_t *lock);
```

- Checking status

- ```
spin_is_locked(spinlock_t *lock);
```

# LKM: Spinlock API (interrupts)

- Between bottom halves
  - Use simple APIs
- Between process context and bottom half
  - spin\_lock\_bh***(*rwlock\_t \*lock*);
  - spin\_unlock\_bh***(*rwlock\_t \*lock*);
  - Soft interrupts on CPU with the lock are disabled
- Between IRQ and bottom halves
  - spin\_lock\_irq***(*rwlock\_t \*lock*);
  - spin\_unlock\_irq***(*rwlock\_t \*lock*);
  - Disable all IRQs before locking, restore all on unlock
- spin\_lock\_irqsave***(*spinlock\_t \*lock*, *unsigned long flags*);
- spin\_unlock\_irqrestore***(*spinlock\_t \*lock*, *unsigned long flags*);
- Save IRQ state before locking, restore exactly on unlock

# LKM: kthread Exercise #3

- Refer **mod8** directory
  - **mod85.c** contains code for spinlock protected resource
    - Again, we share an integer and a char buffer b/w 2 kthreads
    - But we protect using spinlocks
      - Remember we cannot sleep when we have the lock now!
      - *Try sleeping and see what happens!*
    - Compile and load the module on BBB
    - Observe **dmesg** output
      - What do you see as far as kthread-1's prints are concerned?

# LKM: Read-write Spinlock

- Let's assume a shared resource
  - With multiple readers
  - And 1 single writer
- If we use a spinlock
  - Everybody waits
    - Write has to be **exclusive**
    - But what about read?
  - Readers don't really change the resource
    - So multiple readers can actually read at the same time
  - Thus:
    - Write and read have to be mutually exclusive
    - Writes have to be mutually exclusive
    - Any number of reads can be together!
- This situation is handled by a different primitive
  - Called **Read Write Spinlock (RWLock)**

# LKM: RWLock API (simple)

- Creation

- Static

- `DEFINE_RWLOCK(my_rwlock);`

- Dynamic

- `rwlock_t my_rwlock;`

- `rw_lock_init(&my_rwlock);`

- Write Locking / Unlocking

- `write_lock(rwlock_t *lock);`

- `write_unlock(rwlock_t *lock);`

- Read Locking / Unlocking

- `read_lock(rwlock_t *lock);`

- `read_unlock(rwlock_t *lock);`

# LKM: RWLock API (interrupts)

- Between bottom halves

- Use simple APIs

- Between process context and bottom half

***write\_lock\_bh***(rwlock\_t \*lock);

***write\_unlock\_bh***(rwlock\_t \*lock);

- Soft interrupts on CPU with the lock are disabled

***read\_lock\_bh***(rwlock\_t \*lock);

***read\_unlock\_bh***(rwlock\_t \*lock);

- Between IRQ and bottom halves

***write\_lock\_irq***(rwlock\_t \*lock);

***write\_unlock\_irq***(rwlock\_t \*lock);

- Disable all IRQs before locking, restore all on unlock

***read\_lock\_irq***(rwlock\_t \*lock);

***read\_unlock\_irq***(rwlock\_t \*lock);

***write\_lock\_irqsave***(rwlock\_t \*lock, unsigned long flags);

***write\_unlock\_irqrestore***(rwlock\_t \*lock, unsigned long flags);

***read\_lock\_irqsave***(rwlock\_t \*lock, unsigned long flags);

***read\_unlock\_irqrestore***(rwlock\_t \*lock, unsigned long flags);

- Save IRQ state before locking, restore exactly on unlock

# LKM: kthread Exercise #4

- Refer ***mod8*** directory
  - ***mod86.c*** contains code for rwlock protected resource
    - Again, we share an integer b/w 1 write and 2 read threads
    - But we protect using read write locks
    - Compile and load the module on BBB
    - Observe ***dmesg*** output
      - What do you see as far as wthread's prints are concerned?
      - What do you see as far as rthreads' prints are concerned?
    - Change the relative sleep time between read and write
      - Note your observations



# LKM: Sequential Lock

- The issue with the Read Write Lock (RWLock)
  - It favours readers
  - Over writers
- Multiple readers
  - Can starve writers
  - Make writers wait for long times
- What if we wanted to favour writers?
- This situation is handled by a different primitive
  - Called **Sequential Lock (SeqLock)**

# LKM: SeqLock API (simple)

- Creation

```
seqlock_t my_seqlock;  
seq_lock_init(&my_seqlock);
```

- Only the Writer takes the SeqLock

```
write_seqlock(seqlock_t *lock);  
write_sequnlock(seqlock_t *lock);
```

- Reader does not take a lock

- Rather it does a sequence begin – which yields an int (seq\_no)  
*unsigned int seq\_no* = **read\_seqbegin**(*seqlock\_t \*lock*);
- Then it retries the sequence – comparing the earlier seq\_no  
*int* **read\_seqretry**(*seqlock\_t \*lock*, *unsigned int seq\_no*);

# LKM: kthread Exercise #5

- Refer ***mod8*** directory
  - ***mod87.c*** contains code for seqlock protected resource
    - Again, we share an integer b/w 1 write and 1 read thread
    - But we protect using seq locks
    - Compile and load the module on BBB
    - Observe ***dmesg*** output
      - What do you see as far as wthread's prints are concerned?
      - What do you see as far as rthreads' prints are concerned?
    - Change the relative sleep time between read and write
      - Note your observations

THANK YOU!